



AI for Research in Biology

EQ2461 seminar course — Feb 5th, 2025

Antoine Honoré (Postdoc)

Advisor: Prof. Ming Xiao



Outline

Introduction

- Research in Biology
- Cell Functions

Variant effect prediction

- Proteins
- Evolutionary constrain assumption
- Data
- Predictors
- More predictors, based on Transformers

Future perspectives

Conclusions



Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something



Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something



Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something



Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something



Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something



Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something



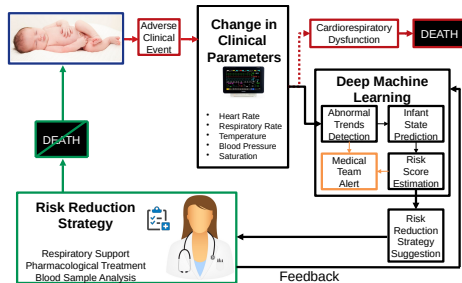
Introduction: Research in Biology

- ▶ Study of living organisms and their interactions with the environment
- ▶ Aim: understanding of the principles and processes that govern life at different levels, e.g. molecular, or generally at organism level:
 - ▶ (Molecular biology) What is the molecular basis for cell function ?
 - ▶ (Physiology) How did specific chemical/physical functions emerge in a biological system ?
- ▶ Research in biology can be either
 - ▶ Basic: Discovering the determinants of functions in living organisms
 - ▶ Applied: Leveraging known processes to achieve something

Introduction: Research in Biology

Human physiology

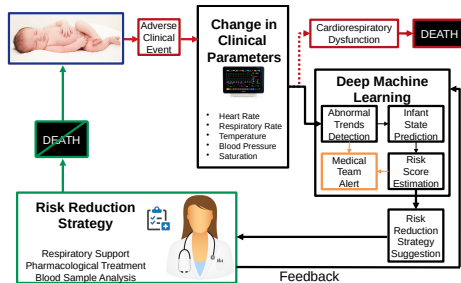
- Basic research: How do infections affect the internal stability of the body ? [1]
- Applied research: Is it possible to leverage the observation of physiological patterns to predict the onset of an infection ? [2]



Introduction: Research in Biology

Human physiology


- ▶ Basic research: How do infections affect the internal stability of the body ? [1]
- ▶ Applied research: Is it possible to leverage the observation of physiological patterns to predict the onset of an infection ? [2]



Introduction: Research in Biology

Human physiology


- ▶ Research: conclusions are limited in part by the ability to collect sufficiently large datasets to account for all sources of variance in the problem
- ▶ Recent techniques allow to collect (and store) large databases, e.g. :
 - ▶ Human physiology : observed with a multitude of high-throughput sensors
 - ▶ Cell biology : variant-cell function maps are sampled abundantly & efficiently with deep mutational scanning experiments



Introduction: Research in Biology

Human physiology


- ▶ Research: conclusions are limited in part by the ability to collect sufficiently large datasets to account for all sources of variance in the problem
- ▶ Recent techniques allow to collect (and store) large databases, e.g. :
 - ▶ Human physiology : observed with a multitude of high-throughput sensors
 - ▶ Cell biology : variant-cell function maps are sampled abundantly & efficiently with deep mutational scanning experiments



Introduction: Research in Biology

Human physiology


- ▶ Research: conclusions are limited in part by the ability to collect sufficiently large datasets to account for all sources of variance in the problem
- ▶ Recent techniques allow to collect (and store) large databases, e.g. :
 - ▶ Human physiology : observed with a multitude of high-throughput sensors
 - ▶ Cell biology : variant-cell function maps are sampled abundantly & efficiently with deep mutational scanning experiments



Introduction: Research in Biology


Human physiology

- ▶ Research: conclusions are limited in part by the ability to collect sufficiently large datasets to account for all sources of variance in the problem
- ▶ Recent techniques allow to collect (and store) large databases, e.g. :
 - ▶ Human physiology : observed with a multitude of high-throughput sensors
 - ▶ Cell biology : variant-cell function maps are sampled abundantly & efficiently with deep mutational scanning experiments



Introduction: Research in Biology

- ▶ AI systems (= models learnt on data) can be used to model large datasets with many interacting variables
- ▶ This allows to make inference in networks such as:




Legend

- ▶ **M**: Medication
- ▶ **L**: Laboratory values (Blood gas analysis)
- ▶ **D**: Demographics (including age)
- ▶ **V**: Mechanical Ventilation
- ▶ **Y**: Clinical condition
- ▶ **X**: Vital signs
- ▶ **G**: Genetic background

- ▶ Instead of being limited to simpler networks such as : $Y \rightarrow X$.

Introduction: Research in Biology

- ▶ AI systems (= models learnt on data) can be used to model large datasets with many interacting variables
- ▶ This allows to make inference in networks such as:



Legend


- ▶ **M**: Medication
- ▶ **L**: Laboratory values (Blood gas analysis)
- ▶ **D**: Demographics (including age)
- ▶ **V**: Mechanical Ventilation
- ▶ **Y**: Clinical condition
- ▶ **X**: Vital signs
- ▶ **G**: Genetic background

- ▶ Instead of being limited to simpler networks such as : $Y \rightarrow X$.

Introduction: Research in Biology

Cell Biology


- Basic research: How do variations (mutations) in the genome affect the function of certain proteins in cells ? [3]
- Applied research: From a variant → cell function prediction tool, can we predict the effect of a chemotherapy drug on a cancer patient ? [4]



Introduction: Research in Biology

Cell Biology

- Basic research: How do variations (mutations) in the genome affect the function of certain proteins in cells ? [3]
- Applied research: From a variant → cell function prediction tool, can we predict the effect of a chemotherapy drug on a cancer patient ? [4]





Introduction: Cell Functions

Introduction: Cell Functions

- ▶ Random genetic variants occur when cells duplicate.
- ▶ Variants may or may not alter the function of the cell.
- ▶ The functions in cells are determined by proteins:




Figure: Each point is related to the 2d-projection of the proteome of a cell for an organism. Each color represents a group of cells with similar function.

Introduction: Cell Functions

- ▶ Random genetic variants occur when cells duplicate.
- ▶ Variants may or may not alter the function of the cell.
- ▶ The functions in cells are determined by proteins:




Figure: Each point is related to the 2d-projection of the proteome of a cell for an organism. Each color represents a group of cells with similar function.

Introduction: Cell Functions

- ▶ Random genetic variants occur when cells duplicate.
- ▶ Variants may or may not alter the function of the cell.
- ▶ The functions in cells are determined by proteins:





Figure: Each point is related to the 2d-projection of the proteome of a cell for an organism. Each color represents a group of cells with similar function.



Variant effect prediction: Proteins



Variant effect prediction: Proteins

In cells

- ▶ Sequences of elementary molecules: amino-acids (AA)
 - ▶ There are $d = 20$ amino-acids (in humans), each with p specific bio-physical properties, e.g. mass, volume.
 - ▶ The length of a protein is typically $L = 300$ AAs (can vary: 10-2000)
- ▶ Determined by their sequence of amino-acid:
PPGPTPLPVIGNILQIGIK...
- ▶ Fold into different 3D-structures because of the bio-physical properties of AAs.



Variant effect prediction: Proteins

In cells

- ▶ Sequences of elementary molecules: amino-acids (AA)
 - ▶ There are $d = 20$ amino-acids (in humans), each with p specific bio-physical properties, e.g. mass, volume.
 - ▶ The length of a protein is typically $L = 300$ AAs (can vary: 10-2000)
- ▶ Determined by their sequence of amino-acid:

PPGPTPLPVIGNILQIGIK...

- ▶ Fold into different 3D-structures because of the bio-physical properties of AAs.

Variant effect prediction: Proteins


In cells

- ▶ Sequences of elementary molecules: amino-acids (AA)
 - ▶ There are $d = 20$ amino-acids (in humans), each with p specific bio-physical properties, e.g. mass, volume.
 - ▶ The length of a protein is typically $L = 300$ AAs (can vary: 10-2000)
- ▶ Determined by their sequence of amino-acid:
PPGPTPLPVIGNILQIGIK...
- ▶ Fold into different 3D-structures because of the bio-physical properties of AAs.

avGFP



Bgl3



GB1



Variant effect prediction: Proteins

Encoding

Proteins can be encoded as

- ① sequences of one-hot vectors in $\{0, 1\}^d$, i.e.
 $\mathbf{x} \in \{0, 1\}^{L \times d}$, (using the alphabet of AAs).
- ② sequences of vectors in \mathbb{R}^p , i.e. $\underline{\mathbf{x}} \in \mathbb{R}^{L \times p}$, (using p bio-physical properties).
- ③ graph signals: $(\underline{\mathbf{x}}, \mathcal{E})$, where there is an edge (collected in the set \mathcal{E}) between two amino-acids (the nodes) if the amino-acids are close in 3D space.

Variant effect prediction: Proteins

Encoding

Proteins can be encoded as


- ① sequences of one-hot vectors in $\{0, 1\}^d$, i.e.
 $\mathbf{x} \in \{0, 1\}^{L \times d}$, (using the alphabet of AAs).
- ② sequences of vectors in \mathbb{R}^p , i.e. $\mathbf{x} \in \mathbb{R}^{L \times p}$, (using p bio-physical properties).
- ③ graph signals: $(\mathbf{x}, \mathcal{E})$, where there is an edge (collected in the set \mathcal{E}) between two amino-acids (the nodes) if the amino-acids are close in 3D space.

Variant effect prediction: Proteins

Encoding

Proteins can be encoded as

- ① sequences of one-hot vectors in $\{0, 1\}^d$, i.e. $\mathbf{x} \in \{0, 1\}^{L \times d}$, (using the alphabet of AAs).
- ② sequences of vectors in \mathbb{R}^p , i.e. $\mathbf{x} \in \mathbb{R}^{L \times p}$, (using p bio-physical properties).
- ③ graph signals: $(\mathbf{x}, \mathcal{E})$, where there is an edge (collected in the set \mathcal{E}) between two amino-acids (the nodes) if the amino-acids are close in 3D space.



Variant effect prediction: Proteins

Why predicting protein variant effects ?

- ▶ In our project we focus on drug transporter proteins:
 - ▶ Membrane proteins involved in drug absorption, distribution, metabolism and excretion (ADME) [5].
- ▶ Loss-of-function variants are directly related to an individual's ability to respond to a drug treatment, e.g. chemotherapy.
- ▶ Being able to predict the effect of variants on these proteins can help select personalized patient treatments.



Variant effect prediction: Proteins

Why predicting protein variant effects ?

- ▶ In our project we focus on drug transporter proteins:
 - ▶ Membrane proteins involved in drug absorption, distribution, metabolism and excretion (ADME) [5].
- ▶ Loss-of-function variants are directly related to an individual's ability to respond to a drug treatment, e.g. chemotherapy.
- ▶ Being able to predict the effect of variants on these proteins can help select personalized patient treatments.



Variant effect prediction: Proteins

Why predicting protein variant effects ?

- ▶ In our project we focus on drug transporter proteins:
 - ▶ Membrane proteins involved in drug absorption, distribution, metabolism and excretion (ADME) [5].
- ▶ Loss-of-function variants are directly related to an individual's ability to respond to a drug treatment, e.g. chemotherapy.
- ▶ Being able to predict the effect of variants on these proteins can help select personalized patient treatments.

Variant effect prediction: Proteins

Why predicting protein variant effects ?

- ▶ In our project we focus on drug transporter proteins:
 - ▶ Membrane proteins involved in drug absorption, distribution, metabolism and excretion (ADME) [5].
- ▶ Loss-of-function variants are directly related to an individual's ability to respond to a drug treatment, e.g. chemotherapy.
- ▶ Being able to predict the effect of variants on these proteins can help select personalized patient treatments.



Variant effect prediction: Evolutionary constrain assumption



Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of positional information content

- ▶ Given N observations of a certain protein, e.g. across species:
 - ▶ Positions where many observations have the same amino-acids are functionally important
 - ▶ Why? because preserved by evolution.
 - ▶ Variants at these positions are considered deleterious
 - ▶ Quantified by the amount of self-information at these positions [6]



Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of positional information content

- ▶ Given N observations of a certain protein, e.g. across species:
 - ▶ Positions where many observations have the same amino-acids are functionally important
 - ▶ Why? because preserved by evolution.
 - ▶ Variants at these positions are considered deleterious
 - ▶ Quantified by the amount of self-information at these positions [6]

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of positional information content

- ▶ Given N observations of a certain protein, e.g. across species:
 - ▶ Positions where many observations have the same amino-acids are functionally important
 - ▶ Why? because preserved by evolution.
 - ▶ Variants at these positions are considered deleterious
 - ▶ Quantified by the amount of self-information at these positions [6]



A sequence logo. From [7].

x-axis: Position in a sequence of AA

y-axis:

- ▶ At each position i , $f_{i,j}$ the frequency of letter j
 - ▶ The entropy: $H_i = -\sum_{j=1}^4 f_{i,j} \log_2 f_{i,j}$
 - ▶ Maximum height = maximum self-information

$$I_{max} = -\log_2(1/4) = 2 \text{ bits}$$

- ▶ Letter j height

$$(I_{max} - H_i) f_{i,j}$$

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of positional information content

- ▶ Given N observations of a certain protein, e.g. across species:
 - ▶ Positions where many observations have the same amino-acids are functionally important
 - ▶ Why? because preserved by evolution.
 - ▶ Variants at these positions are considered deleterious
 - ▶ Quantified by the amount of self-information at these positions [6]



A sequence logo. From [7].

x-axis: Position in a sequence of AA

y-axis:

- ▶ At each position i , $f_{i,j}$ the frequency of letter j
 - ▶ The entropy: $H_i = -\sum_{j=1}^4 f_{i,j} \log_2 f_{i,j}$
 - ▶ Maximum height = maximum self-information

$$I_{max} = -\log_2(1/4) = 2 \text{ bits}$$

- ▶ Letter j height

$$(I_{max} - H_i) f_{i,j}$$

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of distributions over sequence space

► Another way to phrase it, found in [8]:

- ① The distribution of protein sequences observed in nature is the result of billions of evolutionary experiments
- ② Unfit variants were selected out
- ③ By learning a distribution over these sequences we implicitly capture the biochemical constraints that characterize fit variants

The degree of fitness of a variant can be quantified as a log-likelihood ratio:

$$\ln \frac{p(\underline{x}^{(v)}|\theta)}{p(\underline{x}^{(wt)}|\theta)}, \quad (1)$$

- $\underline{x}^{(v)}$ (resp. $\underline{x}^{(wt)}$): the encoded sequence with (resp. without) the variant.
- θ : parameter set, learnt from sequences found in nature only !

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of distributions over sequence space

► Another way to phrase it, found in [8]:

- ① The distribution of protein sequences observed in nature is the result of billions of evolutionary experiments
- ② Unfit variants were selected out
- ③ By learning a distribution over these sequences we implicitly capture the biochemical constraints that characterize fit variants

The degree of fitness of a variant can be quantified as a log-likelihood ratio:

$$\ln \frac{p(\underline{x}^{(v)}|\theta)}{p(\underline{x}^{(wt)}|\theta)}, \quad (1)$$

- $\underline{x}^{(v)}$ (resp. $\underline{x}^{(wt)}$): the encoded sequence with (resp. without) the variant.
- θ : parameter set, learnt from sequences found in nature only !

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of distributions over sequence space

► Another way to phrase it, found in [8]:

- ① The distribution of protein sequences observed in nature is the result of billions of evolutionary experiments
- ② Unfit variants were selected out
- ③ By learning a distribution over these sequences we implicitly capture the biochemical constraints that characterize fit variants

The degree of fitness of a variant can be quantified as a log-likelihood ratio:

$$\ln \frac{p(\underline{x}^{(v)}|\theta)}{p(\underline{x}^{(wt)}|\theta)}, \quad (1)$$

- $\underline{x}^{(v)}$ (resp. $\underline{x}^{(wt)}$): the encoded sequence with (resp. without) the variant.
- θ : parameter set, learnt from sequences found in nature only !

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of distributions over sequence space

► Another way to phrase it, found in [8]:

- ① The distribution of protein sequences observed in nature is the result of billions of evolutionary experiments
- ② Unfit variants were selected out
- ③ By learning a distribution over these sequences we implicitly capture the biochemical constraints that characterize fit variants

The degree of fitness of a variant can be quantified as a log-likelihood ratio:

$$\ln \frac{p(\underline{x}^{(v)}|\theta)}{p(\underline{x}^{(wt)}|\theta)}, \quad (1)$$

- $\underline{x}^{(v)}$ (resp. $\underline{x}^{(wt)}$): the encoded sequence with (resp. without) the variant.
- θ : parameter set, learnt from sequences found in nature only !

Variant effect prediction: Evolutionary constrain assumption

Formulated in terms of distributions over sequence space

► Another way to phrase it, found in [8]:

- ① The distribution of protein sequences observed in nature is the result of billions of evolutionary experiments
- ② Unfit variants were selected out
- ③ By learning a distribution over these sequences we implicitly capture the biochemical constraints that characterize fit variants

The degree of fitness of a variant can be quantified as a log-likelihood ratio:

$$\ln \frac{p(\underline{\mathbf{x}}^{(v)}|\theta)}{p(\underline{\mathbf{x}}^{(wt)}|\theta)}, \quad (1)$$

- $\underline{\mathbf{x}}^{(v)}$ (resp. $\underline{\mathbf{x}}^{(wt)}$): the encoded sequence with (resp. without) the variant.
- θ : parameter set, learnt from sequences found in nature only !

Variant effect prediction: Data

Multiple Sequence Alignment (MSA)

Observations of a protein in different organisms

- ▶ Aligned sequencing outputs for different observations of a protein:

```
DAMMfly2R : MYLPERTEHQKIERLY-----DSNRVN-----AEPGGGL----
DCP1fly2R : -----MTD-----ECVTRNYGVGIRSPNGSENRRGS-FIMADNTDAK-----GCTPESLVVGG
DRICEfly3R : MDATNNGESADQVGIRVGN-----PEQPNDHTDALGSV-GSGGAGSSGLVAGSSHPY-----GSGAIGQLANG
DECAYfly3R : MDDTDFSLFGQKNKHK-----KDKADATKIA-----HTPTSEL----
DRONCfly3L : MQPPELEIGMPKRHREHIRKNLNILVEWTNYERLAMECVQGGILTVQMLRNTQDLNGK-PFNMDEKDVRVEQHRRLLLKITQRGPTAYNLLINA
STRICafly2 : MGWWSKKSETDRSQPSQELVAQDPRTVQTTSAAETETNTAVQNSTITDNNKQTVAFI-FTTRQTVHTQRALITETTRTRTPSQAELEALFAKI
DREDDPafly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDITDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPbfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDITDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPCfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDITDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
```

- ▶ Large databases contain curated MSAs:
 - ▶ Wild-type: A reference sequence, derived from a MSA
 - ▶ Variants: All other sequences.
- ▶ Some variants are indexed in other databases,
 - ▶ associated with clinical phenotypes: needed to score variant effect predictors.

Output

{X} set of sequences to train both positional information-based, and distribution-based models

Variant effect prediction: Data

Multiple Sequence Alignment (MSA)

Observations of a protein in different organisms

- Aligned sequencing outputs for different observations of a protein:

```
DAMMfly2R : MYLPERTEHQKIERLY-----DSNRVN-----AEPGQGL----
DCP1fly2R : -----MTD-----ECVTRNYGVGIRSPNGSENRRGS-FIMADNTDAK-----GCTPESLVVGG
DRICEfly3R : MDATNNGESADQVGIRVGN-----PEQPNHDHTDALGSV-GSGGAGSSGLVAGSSHPY-----GSGAIGQLANG
DECAYfly3R : MDDTDFSLFGQKNKHK-----KDKADATKIA-----HTPTSEL----
DRONCfly3L : MQPPELEIGMPKRHREHIRKNLNILVEWTNYERLAMECVQGGILTVQMLRNTQDLNGK-PFNMDEKDVRVEQHRRLLLKITQRGPTAYNLLINA
STRICafly2 : MGWWSKKSETDRSQPSQELVAQDPRTVQTTSAAETETNTAVQNSTITDNNKQTVAFI-FTTRQTVHTQRALITETTRTRTPSQAELEALFAKI
DREDDPafly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGGLCF-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPbfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGGLCF-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPCfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGGLCF-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
```

- Large databases contain curated MSAs:
 - Wild-type: A reference sequence, derived from a MSA
 - Variants: All other sequences.
- Some variants are indexed in other databases,
 - associated with clinical phenotypes: needed to score variant effect predictors.

Output

{X} set of sequences to train both positional information-based, and distribution-based models

Variant effect prediction: Data

Multiple Sequence Alignment (MSA)

Observations of a protein in different organisms

- Aligned sequencing outputs for different observations of a protein:

```
DAMMfly2R : MYLPERTEHQKIERLY-----DSNRVN-----AEPGQGL----
DCP1fly2R : -----MTD-----ECVTRNYGVGIRSPNGSENRRGS-FIMADNTDAK-----GCTPESLVVGG
DRICEfly3R : MDATNNGESADQVGIRVGN-----PEQPNHDHTDALGSV-GSGGAGSSGLVAGSSHPY-----GSGAIGQLANG
DECAYfly3R : MDDTDFSLFGQKNKHK-----KDKADATKIA-----HTPTSEL----
DRONCfly3L : MQPPELEIGMPKRHREHIRKNLNILVEWTNYERLAMECVQGGILTVQMLRNTQDLNGK-PFNMDEKDVRVEQHRRLLLKITQRGPTAYNLLINA
STRICafly2 : MGWWSKKSETDRSQPSQELVAQDPRTVRQTTSAAETETNTAVQNSTITDNNKQTVAFI-FTROTQVHTQRALITETTRTRTPSQAELEALFAKI
DREDDPafly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPbfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPCfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
```

- Large databases contain curated MSAs:
 - Wild-type: A reference sequence, derived from a MSA
 - Variants: All other sequences.
- Some variants are indexed in other databases,
 - associated with clinical phenotypes: needed to score variant effect predictors.

Output

{X} set of sequences to train both positional information-based, and distribution-based models

Variant effect prediction: Data

Multiple Sequence Alignment (MSA)

Observations of a protein in different organisms

- Aligned sequencing outputs for different observations of a protein:

```
DAMMfly2R : MYLPERTEHQKIERLY-----DSNRVN-----AEPGGGL----
DCP1fly2R : -----MTD-----ECVTRNYGVGIRSPNGSENRRGS-FIMADNTDAK-----GCTPESLVVGG
DRICEfly3R : MDATNNGESADQVGIRVGN-----PEQPNHDHTDALGSV-GSGGAGSSGLVAGSSHPY-----GSGAIGQLANG
DECAYfly3R : MDDTDFSLFGQKNKHK-----KDKADATKIA-----HTPTSEL----
DRONCfly3L : MQPPELEIGMPKRHREHIRKNLNILVEWTNYERLAMECVQGGILTVQMLRNTQDLNGK-PFNMDEKDVRVEQHRRLLLKITQRGPTAYNLLINA
STRICafly2 : MGWWSKKSETDRSQPSQELVAQDPRTVRQTTSAAETETNTAVQNSTITDNNKQTVAFI-TRQTVHTQRALITETTRTRTPSQAELEALFAKI
DREDDPafly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPbfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
DREDDPCfly : MSASAIYRPFPPKVKHFCIFPIAMAGSNLLIHLDTIDQNDLIYVERDMNFAQKVGLGFL-LYGDHSDATYILQKLLAMTRSDFPQSDLLIFFAK
```

- Large databases contain curated MSAs:
 - Wild-type: A reference sequence, derived from a MSA
 - Variants: All other sequences.
- Some variants are indexed in other databases,
 - associated with clinical phenotypes: needed to score variant effect predictors.

Output

{X} set of sequences to train both positional information-based, and distribution-based models



Variant effect prediction: Data

Deep Mutational Scanning

Hard-coding mutations in cells

- ① Create a collection of cells with many single variants
 - ▶ Sequence the collection of cells to get a count of the initial variants (♣)
- ② Expose the cells to an environment
 - ▶ Sequence again and count the remaining variants (♦)
- ③ Compare ♦ and ♣ to get a cell variant "survival score"

Output

$\{(X, y)\}$: X a varied sequence and $y \in \mathbb{R}$ a regression target !

Variant effect prediction: Data

Deep Mutational Scanning

Hard-coding mutations in cells

- ① Create a collection of cells with many single variants
 - ▶ Sequence the collection of cells to get a count of the initial variants (♣)
- ② Expose the cells to an environment
 - ▶ Sequence again and count the remaining variants (♠)
- ③ Compare ♠ and ♣ to get a cell variant "survival score"

Output

$\{(X, y)\}$: X a varied sequence and $y \in \mathbb{R}$ a regression target !




Figure: Position coverage of many experiments

Variant effect prediction: Data

Deep Mutational Scanning

Hard-coding mutations in cells

- ① Create a collection of cells with many single variants
 - ▶ Sequence the collection of cells to get a count of the initial variants (♣)
- ② Expose the cells to an environment
 - ▶ Sequence again and count the remaining variants (♦)
- ③ Compare ♦ and ♣ to get a cell variant "survival score"

Output

$\{(X, y)\}$: X a varied sequence and $y \in \mathbb{R}$ a regression target !




Figure: Position coverage of many experiments

Variant effect prediction: Data

Deep Mutational Scanning

Hard-coding mutations in cells

- ① Create a collection of cells with many single variants
 - ▶ Sequence the collection of cells to get a count of the initial variants (♣)
- ② Expose the cells to an environment
 - ▶ Sequence again and count the remaining variants (♦)
- ③ Compare ♦ and ♣ to get a cell variant "survival score"

Output

$\{(X, y)\}$: X a varied sequence and $y \in \mathbb{R}$ a regression target !




Figure: Position coverage of many experiments




Figure: Density of output scores for 1 experiment

Variant effect prediction: Data

Deep Mutational Scanning

Hard-coding mutations in cells

- ① Create a collection of cells with many single variants
 - ▶ Sequence the collection of cells to get a count of the initial variants (♣)
- ② Expose the cells to an environment
 - ▶ Sequence again and count the remaining variants (♦)
- ③ Compare ♦ and ♣ to get a cell variant "survival score"

Output

$\{(X, y)\}$: X a varied sequence and $y \in \mathbb{R}$ a regression target !




Figure: Position coverage of many experiments




Figure: Density of output scores for 1 experiment

Variant effect prediction: Predictors

Benchmarking

Models are evaluated on variants used in a DMS experiment, and referenced in a clinical database.

- ① x-axis: AUC (Area under the receiver-operating curve)
 - Computed against clinical observation
- ② y-axis: Correlation between the rank predicted and the rank in the DMS experiment.

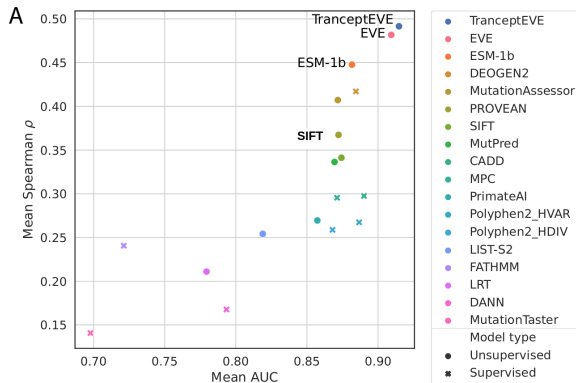


Figure from [9].

Variant effect prediction: Predictors

Benchmarking

Models are evaluated on variants used in a DMS experiment, and referenced in a clinical database.

- ① x-axis: AUC (Area under the receiver-operating curve)
 - Computed against clinical observation
- ② y-axis: Correlation between the rank predicted and the rank in the DMS experiment.

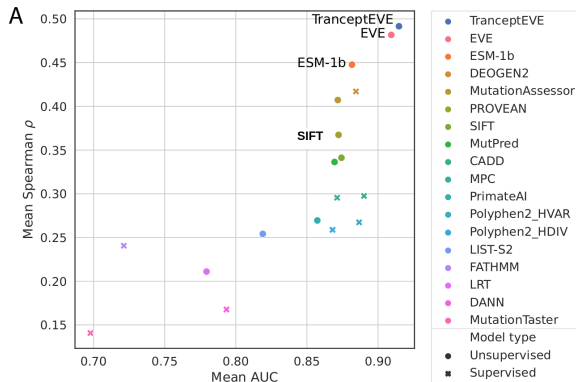


Figure from [9].

Variant effect prediction: Predictors

Benchmarking

Models are evaluated on variants used in a DMS experiment, and referenced in a clinical database.

- ① x-axis: AUC (Area under the receiver-operating curve)
 - Computed against clinical observation
- ② y-axis: Correlation between the rank predicted and the rank in the DMS experiment.




Figure from [9].

Predictors: SIFT

SIFT: Sorting Intolerant From Tolerant

- ① Given an input WT sequence and a variant
- ② Query databases to find similar aligned sequences
- ③ Computes Shannon entropy at each position [10]
- ④ Sum to high entropy \Rightarrow likely benign variant.

Limitations

- ① The prediction depends upon the similar sequences that were queried
- ② Only the position of the variant is taken into account, not the impact on the relation to neighbor AA




Figure from [9].

Predictors: SIFT

SIFT: Sorting Intolerant From Tolerant

- ① Given an input WT sequence and a variant
- ② Query databases to find similar aligned sequences
- ③ Computes Shannon entropy at each position [10]
- ④ Sum to high entropy \Rightarrow likely benign variant.

Limitations

- ① The prediction depends upon the similar sequences that were queried
- ② Only the position of the variant is taken into account, not the impact on the relation to neighbor

AA




Figure from [9].

Predictors: SIFT

SIFT: Sorting Intolerant From Tolerant

- ① Given an input WT sequence and a variant
- ② Query databases to find similar aligned sequences
- ③ Computes Shannon entropy at each position [10]
- ④ Sum to high entropy \Rightarrow likely benign variant.

Limitations

- ① The prediction depends upon the similar sequences that were queried
- ② Only the position of the variant is taken into account, not the impact on the relation to neighbor

AA




Figure from [9].

Predictors: SIFT

SIFT: Sorting Intolerant From Tolerant

- ① Given an input WT sequence and a variant
- ② Query databases to find similar aligned sequences
- ③ Computes Shannon entropy at each position [10]
- ④ Sum to high entropy \implies likely benign variant.

Limitations

- ① The prediction depends upon the similar sequences that were queried
 - ② Only the position of the variant is taken into account, not the impact on the relation to neighbor
- AA




Figure from [9].

Predictors: Deep Sequence

Deep Sequence


A variational auto-encoder (VAE) [3]

- Learn the parameters of the log-likelihood by maximizing a lower bound

Limitations

- ① Vectorized input !
 - Needs retraining on new proteins
- ② Unimodal distribution in the latent space
 - Limit expressivity of latent vectors
- ③ For inference, the ratio of likelihoods is replaced with a ratio of lower bounds:

$$\ln \frac{p(\underline{x}^v | \theta)}{p(\underline{x}^{wt} | \theta)} \approx \ln \frac{\mathcal{L}_{VLB}(\underline{x}^v | \theta)}{\mathcal{L}_{VLB}(\underline{x}^{wt} | \theta)} \quad (2)$$



Predictors: Deep Sequence

Deep Sequence


A variational auto-encoder (VAE) [3]

- Learn the parameters of the log-likelihood by maximizing a lower bound

Limitations

- ① Vectorized input !
 - Needs retraining on new proteins
- ② Unimodal distribution in the latent space
 - Limit expressivity of latent vectors
- ③ For inference, the ratio of likelihoods is replaced with a ratio of lower bounds:

$$\ln \frac{p(\underline{x}^v | \theta)}{p(\underline{x}^{wt} | \theta)} \approx \ln \frac{\mathcal{L}_{VLB}(\underline{x}^v | \theta)}{\mathcal{L}_{VLB}(\underline{x}^{wt} | \theta)} \quad (2)$$



Predictors: Deep Sequence

Deep Sequence


A variational auto-encoder (VAE) [3]

- Learn the parameters of the log-likelihood by maximizing a lower bound

Limitations

- ① Vectorized input !
 - Needs retraining on new proteins
- ② Unimodal distribution in the latent space
 - Limit expressivity of latent vectors
- ③ For inference, the ratio of likelihoods is replaced with a ratio of lower bounds:

$$\ln \frac{p(\underline{x}^v | \theta)}{p(\underline{x}^{wt} | \theta)} \approx \ln \frac{\mathcal{L}_{VLB}(\underline{x}^v | \theta)}{\mathcal{L}_{VLB}(\underline{x}^{wt} | \theta)} \quad (2)$$



Predictors: Deep Sequence

Deep Sequence


A variational auto-encoder (VAE) [3]

- Learn the parameters of the log-likelihood by maximizing a lower bound

Limitations

- ① Vectorized input !
 - Needs retraining on new proteins
- ② Unimodal distribution in the latent space
 - Limit expressivity of latent vectors
- ③ For inference, the ratio of likelihoods is replaced with a ratio of lower bounds:

$$\ln \frac{p(\underline{x}^v|\theta)}{p(\underline{x}^{wt}|\theta)} \approx \ln \frac{\mathcal{L}_{VLB}(\underline{x}^v|\theta)}{\mathcal{L}_{VLB}(\underline{x}^{wt}|\theta)} \quad (2)$$



Predictors: Deep Sequence

Deep Sequence


A variational auto-encoder (VAE) [3]

- Learn the parameters of the log-likelihood by maximizing a lower bound

Limitations

- ① Vectorized input !
 - Needs retraining on new proteins
- ② Unimodal distribution in the latent space
 - Limit expressivity of latent vectors
- ③ For inference, the ratio of likelihoods is replaced with a ratio of lower bounds:

$$\ln \frac{p(\mathbf{x}^v|\theta)}{p(\mathbf{x}^{wt}|\theta)} \approx \ln \frac{\mathcal{L}_{VLB}(\mathbf{x}^v|\theta)}{\mathcal{L}_{VLB}(\mathbf{x}^{wt}|\theta)} \quad (2)$$





Predictors: Deep Sequence

My current work: addressing the limitations of DeepSequence

- ▶ Allow multimodality in the latent space:
 - ▶ Mixture of Gaussian, instead of Gaussian ?
 - ▶ Did not work, the mixture was not used by the model.
- ▶ Better idea:
 - ▶ Latent space = probability simplex
 - ▶ Replace D_{KL} with entropy of latent vector in the lower bound


$$\text{Recall: } \ln p(\mathbf{x}) = D_{KL}(q_\phi(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) - D_{KL}(q_\phi(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) + \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\ln p_\theta(\mathbf{x}|\mathbf{z})]$$

Predictors: Deep Sequence

My current work: addressing the limitations of DeepSequence

- Less parameters (8-9M) (no need for Bayesian decoder)
- More interpretability,
- Better/similar results!

Model name	Spearmanr
Best Benchmark	0.529 \pm 0.151
EVE (single)	0.432 \pm 0.161
matVAE-MSA	0.428 \pm 0.15
DeepSequence (single)	0.412 \pm 0.149
Site-Independent	0.39 \pm 0.145



More predictors, based on Transformers: ESM

ESM: Evolutionary Scale Modeling

A protein model with 650M parameters [11]


- Given an input varied sequence $\underline{\mathbf{x}}^v = (\mathbf{x}_1, \dots, \mathbf{x}_L)^T$, the statistical model is similar to a large language model :

$$p(\underline{\mathbf{x}}^v | \theta) = \prod_{i=1}^L p(\mathbf{x}_i | \mathbf{x}_{N(i)}, \theta), \quad (3)$$

with $N(i)$ a randomly masked neighborhood for i .

- The prediction score is the ratio

$$\ln \frac{p(\underline{\mathbf{x}}^v | \theta)}{p(\underline{\mathbf{x}}^{wt} | \theta)} \quad (4)$$



More predictors, based on Transformers: ESM

ESM: Evolutionary Scale Modeling

A protein model with 650M parameters [11]


- ① Given an input varied sequence $\underline{\mathbf{x}}^v = (\mathbf{x}_1, \dots, \mathbf{x}_L)^T$, the statistical model is similar to a large language model :

$$p(\underline{\mathbf{x}}^v | \theta) = \prod_{i=1}^L p(\mathbf{x}_i | \mathbf{x}_{N(i)}, \theta), \quad (3)$$

with $N(i)$ a randomly masked neighborhood for i .

- ② The prediction score is the ratio

$$\ln \frac{p(\underline{\mathbf{x}}^v | \theta)}{p(\underline{\mathbf{x}}^{wt} | \theta)} \quad (4)$$



More predictors, based on Transformers: ESM

ESM: Evolutionary Scale Modeling

A protein model with 650M parameters [11]


- Given an input varied sequence $\underline{\mathbf{x}}^v = (\mathbf{x}_1, \dots, \mathbf{x}_L)^T$, the statistical model is similar to a large language model :

$$p(\underline{\mathbf{x}}^v | \theta) = \prod_{i=1}^L p(\mathbf{x}_i | \mathbf{x}_{N(i)}, \theta), \quad (3)$$

with $N(i)$ a randomly masked neighborhood for i .

- The prediction score is the ratio

$$\ln \frac{p(\underline{\mathbf{x}}^v | \theta)}{p(\underline{\mathbf{x}}^{wt} | \theta)} \quad (4)$$



More predictors, based on Transformers: TranceptEVE

TranceptEVE: Tranception + EVE

An auto-regressive protein language model combined with other models [3]

- ▶ Tranception [12]: auto-regressive model, i.e. similar to ESM but the neighborhood is all the previous AAs.

$$p_T(\mathbf{x}_1, \dots, \mathbf{x}_L | \theta) = \prod_{i=1}^L p_T(\mathbf{x}_i | \mathbf{x}_{i-1}, \dots, \mathbf{x}_1, \theta). \quad (5)$$

- ▶ The final log-likelihood is computed as a convex combination of log-likelihood from other models:

$$\log p(\mathbf{x} | \theta) \propto \sum_{i=1}^L (1 - \beta) [(1 - \alpha) \log p_T(\mathbf{x}_i | \mathbf{x}_{<i}) + \alpha \log p_R(\mathbf{x}_i)] + \beta \log p_E(\mathbf{x}_i), \quad (6)$$

- ▶ p_R is the empirical distribution from the MSA, p_E is the output of EVE, (α, β) are hyper-parameters which favor p_T when the MSA is shallow, i.e. unreliable.

More predictors, based on Transformers: TranceptEVE

TranceptEVE: Tranception + EVE

An auto-regressive protein language model combined with other models [3]

- ▶ Tranception [12]: auto-regressive model, i.e. similar to ESM but the neighborhood is all the previous AAs.

$$p_T(\mathbf{x}_1, \dots, \mathbf{x}_L | \theta) = \prod_{i=1}^L p_T(\mathbf{x}_i | \mathbf{x}_{i-1}, \dots, \mathbf{x}_1, \theta). \quad (5)$$

- ▶ The final log-likelihood is computed as a convex combination of log-likelihood from other models:

$$\log p(\mathbf{x} | \theta) \propto \sum_{i=1}^L (1 - \beta) [(1 - \alpha) \log p_T(\mathbf{x}_i | \mathbf{x}_{<i}) + \alpha \log p_R(\mathbf{x}_i)] + \beta \log p_E(\mathbf{x}_i), \quad (6)$$

- ▶ p_R is the empirical distribution from the MSA, p_E is the output of EVE, (α, β) are hyper-parameters which favor p_T when the MSA is shallow, i.e. unreliable.

More predictors, based on Transformers: TranceptEVE

TranceptEVE: Tranception + EVE

An auto-regressive protein language model combined with other models [3]

- ▶ Tranception [12]: auto-regressive model, i.e. similar to ESM but the neighborhood is all the previous AAs.

$$p_T(\mathbf{x}_1, \dots, \mathbf{x}_L | \theta) = \prod_{i=1}^L p_T(\mathbf{x}_i | \mathbf{x}_{i-1}, \dots, \mathbf{x}_1, \theta). \quad (5)$$

- ▶ The final log-likelihood is computed as a convex combination of log-likelihood from other models:

$$\log p(\mathbf{x} | \theta) \propto \sum_{i=1}^L (1 - \beta) [(1 - \alpha) \log p_T(\mathbf{x}_i | \mathbf{x}_{<i}) + \alpha \log p_R(\mathbf{x}_i)] + \beta \log p_E(\mathbf{x}_i), \quad (6)$$

- ▶ p_R is the empirical distribution from the MSA, p_E is the output of EVE, (α, β) are hyper-parameters which favor p_T when the MSA is shallow, i.e. unreliable.

More predictors, based on Transformers: TranceptEVE

TranceptEVE: Tranception + EVE

An auto-regressive protein language model combined with other models [3]

- ▶ Tranception [12]: auto-regressive model, i.e. similar to ESM but the neighborhood is all the previous AAs.

$$p_T(\mathbf{x}_1, \dots, \mathbf{x}_L | \theta) = \prod_{i=1}^L p_T(\mathbf{x}_i | \mathbf{x}_{i-1}, \dots, \mathbf{x}_1, \theta). \quad (5)$$

- ▶ The final log-likelihood is computed as a convex combination of log-likelihood from other models:

$$\log p(\mathbf{x} | \theta) \propto \sum_{i=1}^L (1 - \beta) [(1 - \alpha) \log p_T(\mathbf{x}_i | \mathbf{x}_{<i}) + \alpha \log p_R(\mathbf{x}_i)] + \beta \log p_E(\mathbf{x}_i), \quad (6)$$

- ▶ p_R is the empirical distribution from the MSA, p_E is the output of EVE, (α, β) are hyper-parameters which favor p_T when the MSA is shallow, i.e. unreliable.

Future perspectives

- ① Existing prediction models remain reliant on evolutionary constrain assumption:
 - ▶ Either through the informational or probabilistic formulation, or both.
 - ▶ Problem: Drug transporter proteins were not constrained by evolution (similar to nutriment transport in and out of cells)
- ② Many models do not take into account the 3d structure of proteins:
 - ▶ A few works used graph neural networks, did not show very big improvements over e.g. CNN
- ③ Deep mutational scanning data are promising
 - ▶ Problem:
 - ▶ lack of standardization of experimental methods, i.e. output are not always numerically comparable.
 - ▶ Advantages:
 - ▶ Can provide information on specific phenotype,
 - ▶ Train against quantitative functional scores

Future perspectives

- ① Existing prediction models remain reliant on evolutionary constrain assumption:
 - ▶ Either through the informational or probabilistic formulation, or both.
 - ▶ Problem: Drug transporter proteins were not constrained by evolution (similar to nutriment transport in and out of cells)
- ② Many models do not take into account the 3d structure of proteins:
 - ▶ A few works used graph neural networks, did not show very big improvements over e.g. CNN
- ③ Deep mutational scanning data are promising
 - ▶ Problem:
 - ▶ lack of standardization of experimental methods, i.e. output are not always numerically comparable.
 - ▶ Advantages:
 - ▶ Can provide information on specific phenotype,
 - ▶ Train against quantitative functional scores

Future perspectives

- ① Existing prediction models remain reliant on evolutionary constrain assumption:
 - ▶ Either through the informational or probabilistic formulation, or both.
 - ▶ Problem: Drug transporter proteins were not constrained by evolution (similar to nutriment transport in and out of cells)
- ② Many models do not take into account the 3d structure of proteins:
 - ▶ A few works used graph neural networks, did not show very big improvements over e.g. CNN
- ③ Deep mutational scanning data are promising
 - ▶ Problem:
 - ▶ lack of standardization of experimental methods, i.e. output are not always numerically comparable.
 - ▶ Advantages:
 - ▶ Can provide information on specific phenotype,
 - ▶ Train against quantitative functional scores

Future perspectives

- ① Existing prediction models remain reliant on evolutionary constrain assumption:
 - ▶ Either through the informational or probabilistic formulation, or both.
 - ▶ Problem: Drug transporter proteins were not constrained by evolution (similar to nutriment transport in and out of cells)
- ② Many models do not take into account the 3d structure of proteins:
 - ▶ A few works used graph neural networks, did not show very big improvements over e.g. CNN
- ③ Deep mutational scanning data are promising
 - ▶ Problem:
 - ▶ lack of standardization of experimental methods, i.e. output are not always numerically comparable.
 - ▶ Advantages:
 - ▶ Can provide information on specific phenotype,
 - ▶ Train against quantitative functional scores

Future perspectives

- ① Existing prediction models remain reliant on evolutionary constrain assumption:
 - ▶ Either through the informational or probabilistic formulation, or both.
 - ▶ Problem: Drug transporter proteins were not constrained by evolution (similar to nutriment transport in and out of cells)
- ② Many models do not take into account the 3d structure of proteins:
 - ▶ A few works used graph neural networks, did not show very big improvements over e.g. CNN
- ③ Deep mutational scanning data are promising
 - ▶ Problem:
 - ▶ lack of standardization of experimental methods, i.e. output are not always numerically comparable.
 - ▶ Advantages:
 - ▶ Can provide information on specific phenotype,
 - ▶ Train against quantitative functional scores

Future perspectives

- ① Existing prediction models remain reliant on evolutionary constrain assumption:
 - ▶ Either through the informational or probabilistic formulation, or both.
 - ▶ Problem: Drug transporter proteins were not constrained by evolution (similar to nutriment transport in and out of cells)
- ② Many models do not take into account the 3d structure of proteins:
 - ▶ A few works used graph neural networks, did not show very big improvements over e.g. CNN
- ③ Deep mutational scanning data are promising
 - ▶ Problem:
 - ▶ lack of standardization of experimental methods, i.e. output are not always numerically comparable.
 - ▶ Advantages:
 - ▶ Can provide information on specific phenotype,
 - ▶ Train against quantitative functional scores

Conclusions

- ① We have seen what research in biology focuses on: understanding the processes that govern life
- ② We have seen how AI systems can be used to help research in biology
 - ▶ i.e. have the potential to provide conclusive answers to difficult questions by integrating the large amount of data generated by technological advances
- ③ Certain models are at the forefront of research: graph neural networks, transformers, VAEs
- ④ Engineering efficient analysis pipelines remains important to harness powerful models and large datasets

Thank you !



References I

- [1] V. Siljehav, A. M. Hofstetter, K. Leifsdottir, and E. Herlenius, “Prostaglandin E2 Mediates Cardiorespiratory Disturbances during Infection in Neonates,” vol. 167, no. 6, pp. 1207–1213.e3.
- [2] A. Honoré, D. Forsberg, K. Adolphson, S. Chatterjee, K. Jost, and E. Herlenius, “Vital sign-based detection of sepsis in neonates using machine learning,” vol. n/a, no. n/a.
- [3] A. J. Riesselman, J. B. Ingraham, and D. S. Marks, “Deep generative models of genetic variation capture the effects of mutations,” vol. 15, no. 10, pp. 816–822.
- [4] Y. Zhou, R. Tremmel, E. Schaeffeler, M. Schwab, and V. M. Lauschke, “Challenges and opportunities associated with rare-variant pharmacogenomics,” vol. 43, no. 10, pp. 852–865.
- [5] A. G. Roberts, “The Structure and Mechanism of Drug Transporters,” vol. 2342, pp. 193–234.

References II

- [6] T. D. Schneider, G. D. Stormo, L. Gold, and A. Ehrenfeucht, “Information content of binding sites on nucleotide sequences,” vol. 188, no. 3, pp. 415–431.
- [7] T. D. Schneider and R. M. Stephens, “Sequence logos: A new way to display consensus sequences,” vol. 18, no. 20, pp. 6097–6100.
- [8] P. Notin, L. V. Niekerk, A. W. Kollasch, D. Ritter, Y. Gal, and D. S. Marks, “TranceptEVE: Combining Family-specific and Family-agnostic Models of Protein Sequences for Improved Fitness Prediction.”
- [9] P. Notin, A. W. Kollasch, D. Ritter, L. V. Niekerk, S. Paul, H. Spinner, N. J. Rollins, A. Shaw, R. Orenbuch, R. Weitzman, J. Frazer, M. Dias, D. Franceschi, Y. Gal, and D. S. Marks, “ProteinGym: Large-Scale Benchmarks for Protein Fitness Prediction and Design,”
- [10] P. C. Ng and S. Henikoff, “Predicting Deleterious Amino Acid Substitutions,” vol. 11, no. 5, pp. 863–874.

References III

- [11] A. Rives, J. Meier, T. Sercu, S. Goyal, Z. Lin, J. Liu, D. Guo, M. Ott, C. L. Zitnick, J. Ma, and R. Fergus, “Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences,” vol. 118, no. 15, p. e2016239118.
- [12] P. Notin, M. Dias, J. Frazer, J. M. Hurtado, A. N. Gomez, D. Marks, and Y. Gal, “Tranception: Protein Fitness Prediction with Autoregressive Transformers and Inference-time Retrieval,” in *Proceedings of the 39th International Conference on Machine Learning*, pp. 16990–17017, PMLR.